SUPERCHARGE YOUR AWS EXPERTISE WITH GENERATIVE AI

SR. PARTNER SA NIKO BORODACHUK

AMAZON WEB SERVICES

What is Generative Artificial Intelligence (GenAI)?

- Creates new content and ideas, including conversations, stories, images, videos, and music
- Powered by large models that are pretrained on vast corpuses of data and commonly referred to as foundation models (FMs)

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Generative AI is powered by foundation models

Pretrained on vast amounts of unstructured data

Contain large number of parameters that make them capable of learning complex concepts

Can be applied in a wide range of contexts

Customize FMs using your data for domain specific tasks

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Enhance Customer Experiences

CHATBOTS

VIRTUAL ASSISTANTS

CONVERSATION ANALYTICS

PERSONALIZATION

Boost employee productivity & creativity

CONVERSATIONAL SEARCH

SUMMARIZATION

CONTENT CREATION

CODE GENERATION

DATA TO INSIGHTS

Optimize business processes

DOCUMENT PROCESSING

DATA AUGMENTATION

CYBERSECURITY

PROCESS OPTIMIZATION

What types of foundation models are currently in the market?

Text-to-text

These natural language processing models can summarize text, extract information, respond to questions, and create content such as blogs or product descriptions.

An example is sentence auto-completion.

Text-to-embeddings

These FMs compare user search bar input with index data and connect the dots between the two. The result is more accurate and relevant results.

Multimodal

These emerging foundation models can generate images based on a user's natural language text input.

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Why Build generative AI with AWS

	Customer challenge:	AWS solution:
Flexibility	Customers need a straightforward way to find and access high-performing FMs that are best suited for their business.	AWS offers a wide selection of FMs built by Amazon and top AI startups, including AI21 Labs, Anthropic, and Stability AI.
Secure customization	Customer need their data to stay completely protected, secure, and private.	Using Amazon Bedrock, customers can fine tune models for a particular task without having to annotate large volumes of data (as few as 20 examples is enough). No customer data is used to train the underlying models. All data is encrypted and stays within the customer's virtual private cloud (VPC).
Most cost-effective infrastructure:	To fully use foundation models, customers need the most high-performing, cost-effective infrastructure purpose-built for ML.	AWS offers the best price performance for generative AI with infrastructure powered by AWS designed ML chips and NVIDIA GPUs.
Easiest way to build with FMs	Customers want ease of use. They want to quickly integrate and deploy FMs into their applications and workloads running on AWS.	Using AWS generative AI capabilities, customers don't need to send their data to the model. Instead, they can bring the model to their data using familiar controls and integrations and services such as Amazon SageMaker and Amazon Simple Storage Service (Amazon S3).
Generative AI-powered solutions	Customers are looking for generative AI solutions to help improve productivity while seamlessly interacting with applications and systems.	With generative AI built in, services such as Amazon CodeWhisperer, an AI coding companion, can help customers improve productivity.

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Everything you need to accelerate your generative AI journey

Easiest way to build

with leading foundation models

Differentiate with your data

in a secure and private environment

Increase productivity

with generative AI applications and services

Most performant, low cost

infrastructure to scale generative AI

NOW GENERALLY AVAILABLE

Amazon **Bedrock**

The easiest way to build and scale generative AI applications with foundation models (FMs)

Accelerate development of generative AI applications using FMs through an API, without managing infrastructure

Choose FMs from Amazon, AI21 Labs, Anthropic, Cohere, Meta, and Stability AI to find the right FM for your use case

Privately customize FMs using your organization's data

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Everything you need to accelerate your generative AI journey

Easiest way to build

with leading foundation models

Differentiate with your data

in a secure and private environment

Increase productivity

with generative AI applications and services

Most performant, low cost

infrastructure to scale generative AI

Keeping your data private and secure

None of the customer's data is used to train the underlying model

All data is encrypted at rest and PrivateLink support allows access to Bedrock APIs via customer's VPC endpoints

Customized foundation models and the customer-specific data that trains them remain private

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Everything you need to accelerate your generative AI journey

Easiest way to build

with leading foundation models

Differentiate with your data

in a secure and private environment

Increase productivity

with generative AI applications and services

Most performant, low cost

infrastructure to scale generative AI

GENERALLY AVAILABLE

Amazon **CodeWhisperer**

Build apps faster and more securely with an AI coding companion

Generate code suggestions in real-time

Scan code for hard-to-find vulnerabilities

Flag code that resembles open-source training data or filter by default

FREE FOR INDIVIDUAL TIER

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Everything you need to accelerate your generative AI journey

Easiest way to build

with leading foundation models

Differentiate with your data

in a secure and private environment

Increase productivity

with generative AI applications and services

Most performant, low cost

infrastructure to scale generative AI

Purpose-built accelerators for generative AI

AWS Trainium

Up to 50% savings on training costs over comparable Amazon EC2 instances

AWS Inferentia2

Up to 40% better price performance than comparable Amazon EC2 instances

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon SageMaker JumpStart

ML hub with foundation models, built-in algorithms, and prebuilt ML solutions that you can deploy with just a few clicks

Machine learning hub

Browse through 400+ built-in algorithms with pretrained foundation models, solutions, and example notebooks

Pre-built training and inference scripts

Compatible with SageMaker and configurable with custom dataset

UI as well as API-based

Use the user interface for single click model deployment or API for the Python SDK-based workflow

Notebooks with examples

Jump into notebooks to use selected model with examples to guide you through the entire ML workflow

Share and collaborate within your organization

Share models and notebooks with others within your organization, and allow them to train with their own data or deploy as-is for inferencing

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Write your Badge Exam

GenAI business

Skills Builder assessment

Partner central self register

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

THANK YOU

© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.